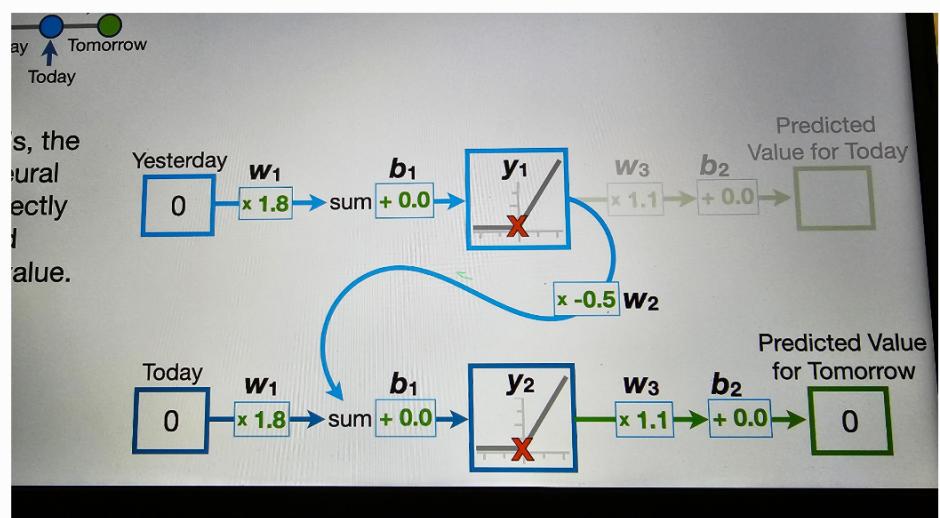
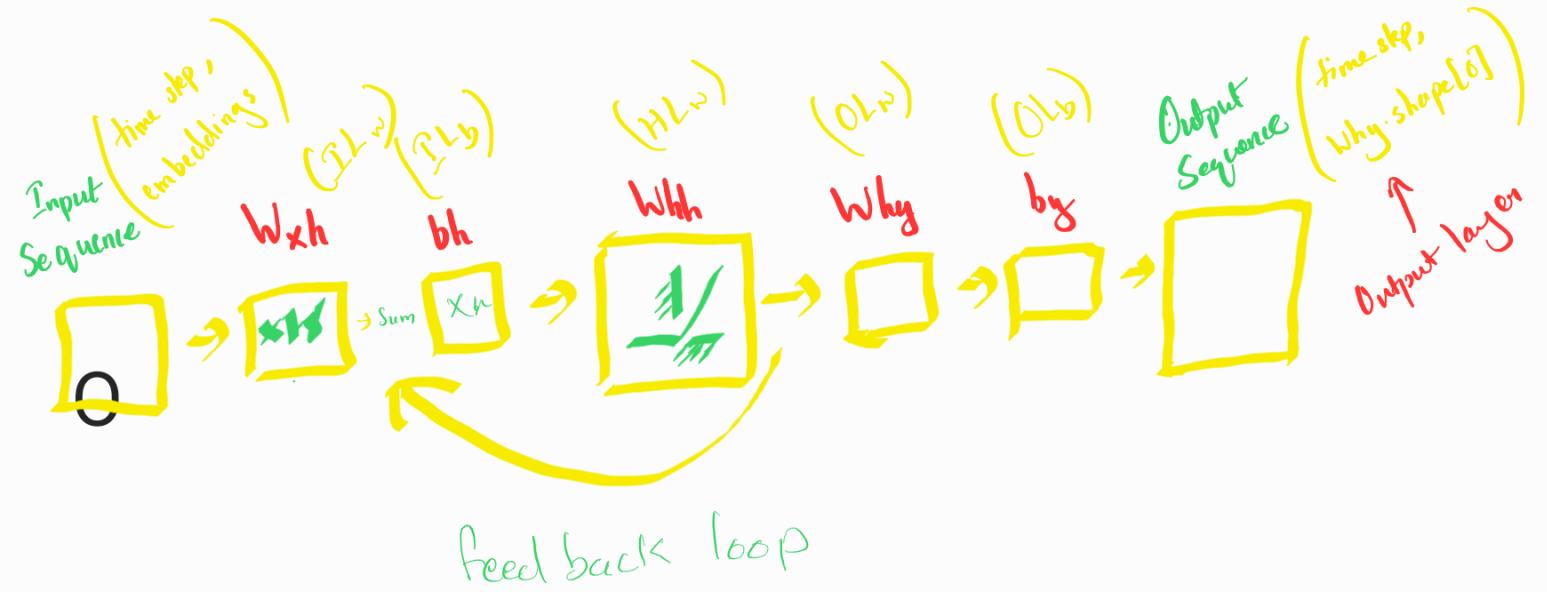


- RNN have weights, biases layers, activation functions
- The difference is that it has a feed back loop.
- The feed back loop makes it possible to use sequential inputs one by one over time



- Instead of using the Input of the 1st for prediction it is given through the feedback loop and makes a prediction with the sequence of Inputs.

- Initialize hidden state $(1, Whh.shape[0])$

- When each input comes in we have to unroll the network. (results in new neural networks in background)

- Looping through the time step
- Finding input at current step
- Compute hidden state at current time step

$$h_s = \text{tf.matmul}(x_t, W_{xh}) + \text{tf.matmul}(h_{t-1}, W_{hh}) + b_h$$

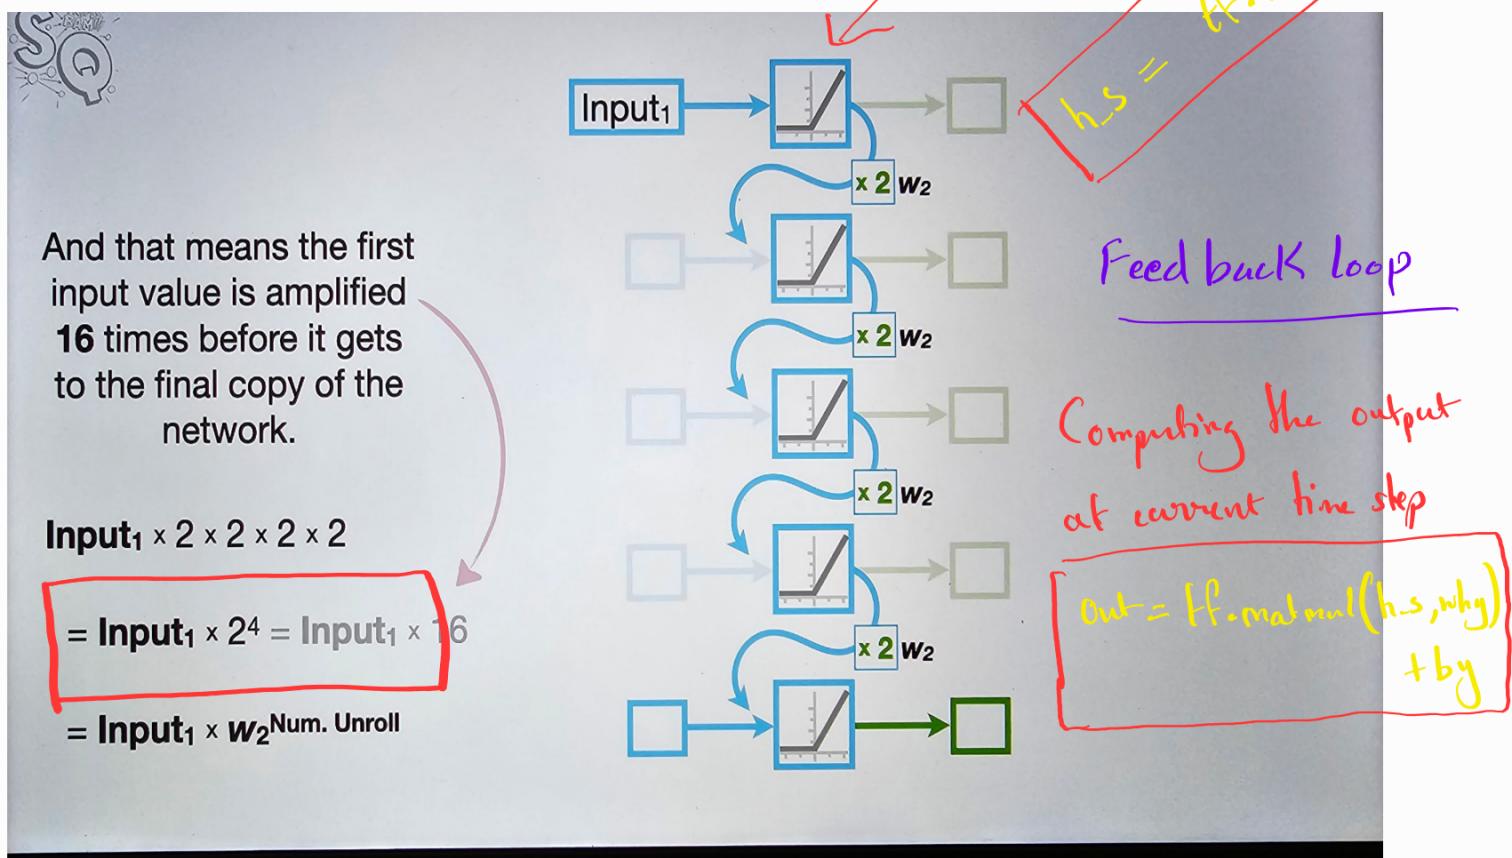
- Even though we unroll the RNN multiple times the weights and biases are shared and there is no new addition of extra weights.

- The more we unroll RNN the more harder it becomes to

Vanishing / Exploding Gradient Descent problem)

- where did the gradient go?

Exploding Problem



- The Input gets multiplied

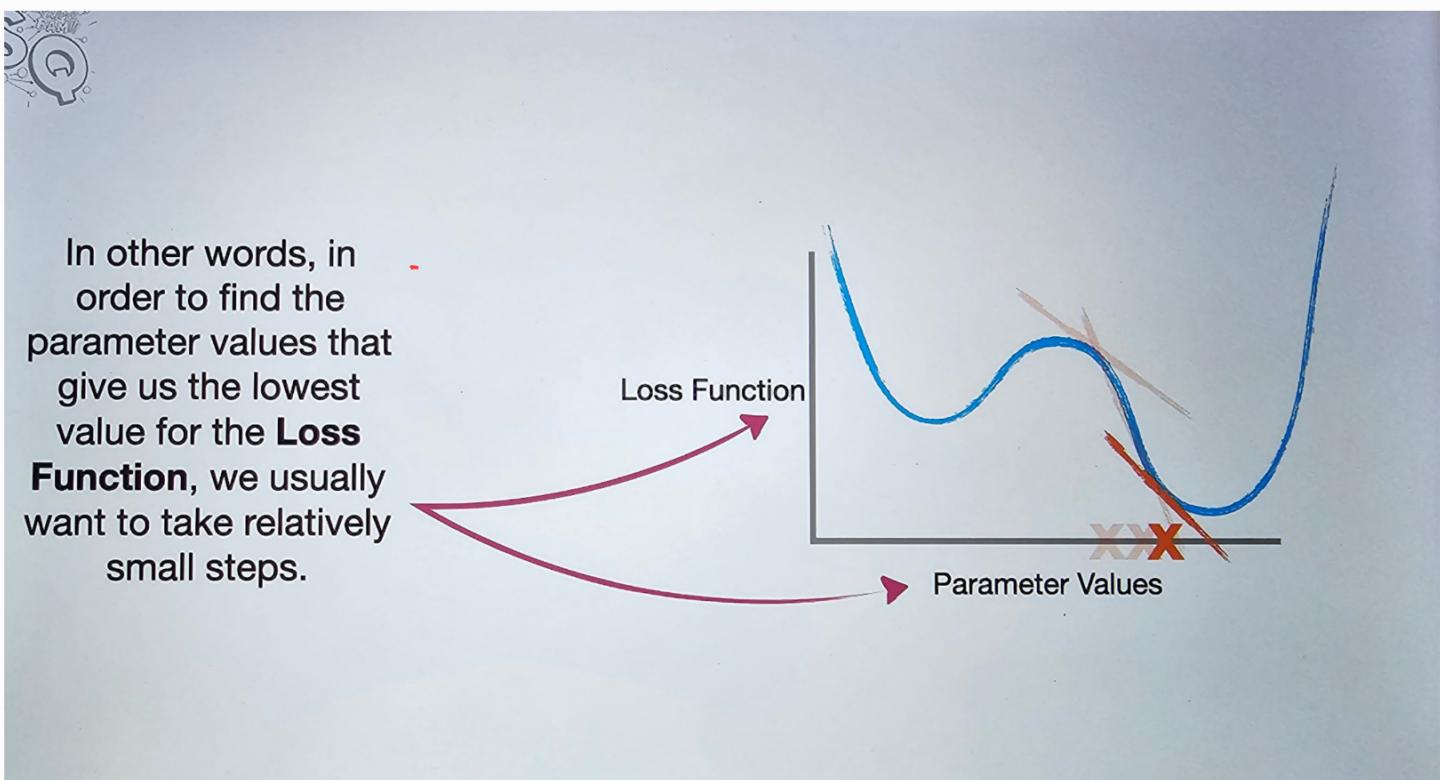
as many times as we

unroll the RNN.

50

- If we have 50 data 1×2

\Rightarrow ! HUGE NUMBER



- It neglects us from taking

small steps and we end up in Exploding GD. We just bounce around.

Vanishing Gradient problem

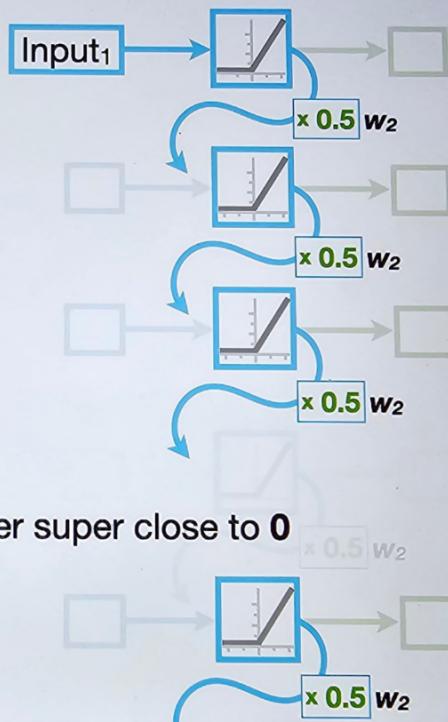
- When we limit w_2 to values < 1
- A solution to avoid Exploding Gradient Descent.



Because this number is **super close to 0**, this is called the **Vanishing Gradient Problem**.

$$\text{Input}_1 \times 0.5^{50} = \text{Input}_1 \times \text{a number super close to 0}$$

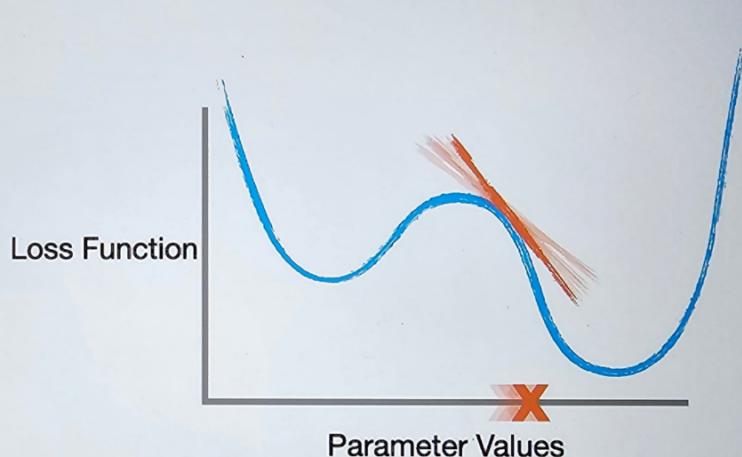
$\text{Input}_1 \times w_2^{\text{Num. Unroll}}$



- Results in Vanishing
GD



And as a result, we end up hitting the maximum number of steps we are allowed to take before we find the optimal value.



- Solution to Exploding /

Vanishing GD is LSTM

(Long - short - term

Memory - Networks

Training a RNN

Parameters W_{xh} , W_{hh} ,

W_y , b_h , b_y ,

input Sequence , targets,
alpha, epochs .

Steps

1) Defining the Optimizer

Optimizer = tf. optimizers. Adam

(learning_rate = alpha)

2) Defining the loss fn

`loss = tf.losses.MeanSquaredError()`

3) Looping Through Epochs -

4) Iterating Through

Each batch of Input

`(Input_sequence.shape[0])`

5) with `tf.GradientTape()` as

`tape`:

6) Finding y_{pred}

Output sequence, $h-s$

$$= \text{rnn-f-p}(\text{w}_{xh}, \text{w}_{hh}, \text{w}_{hy} \\ \text{bh}, \text{by}, \text{batch - input})$$

7) finding loss

$$\text{loss} = \text{loss_fn}(\text{batch_target}, \\ \text{Output Sequence})$$

8) Finding Gradients in the Scope of with GradientTape() as tape:

grads = tape.gradient(loss,
[Wxh, Whh, Why, bh,
by])

9) Applying Optimizer

On Gradients.

Optimizer.apply_gradients

(zip(grads, [Wxh, Whh, Why,
bh, by]))

10) returning Wxh, Whh,
Why, bh, by

